

Affinity & KOS

Jalal Kawash

Linux affinity sys calls

http://linux.die.net/man/2/sched_getaffinity

```
int sched_setaffinity(pid_t pid, size_t cpusetsize, cpu_set_t *mask);  
int sched_getaffinity(pid_t pid, size_t cpusetsize, cpu_set_t *mask);
```

```
int sched_setaffinity(pid_t pid, size_t  
cpusetsize, cpu_set_t *mask);
```

- “sets the CPU affinity mask of the process whose ID is *pid* to the value specified by *mask*. If *pid* is zero, then the calling process is used. The argument *cpusetsize* is the length (in bytes) of the data pointed to by *mask*. Normally this argument would be specified as *sizeof(cpu_set_t)*.
- If the process specified by *pid* is not currently running on one of the CPUs specified in *mask*, then that process is migrated to one of the CPUs specified in *mask*. “

```
int sched_getaffinity(pid_t pid, size_t  
cpusetsize, cpu_set_t *mask);
```

- “writes the affinity mask of the process whose ID is *pid* into the *cpu_set_t* structure pointed to by *mask*. The *cpusetsize* argument specifies the size (in bytes) of *mask*. If *pid* is zero, then the mask of the calling process is returned.
- **Return Value**
- On success, **sched_setaffinity()** and **sched_getaffinity()** return 0. On error, -1 is returned, and *errno* is set appropriately. “

Errors

Errors

EFAULT

A supplied memory address was invalid.

EINVAL

The affinity bit mask *mask* contains no processors that are currently physically on the system and permitted to the

EINVAL

(**sched_getaffinity()** and, in kernels before 2.6.9, **sched_setaffinity()**) *cpusetsize* is smaller than the size of the

EPERM

(**sched_setaffinity()**) The calling process does not have appropriate privileges. The caller needs an effective user

ESRCH

The process whose ID is *pid* could not be found.

cpu_set_t

- `cpu_set_t` is already defined as a 64 bit unsigned integer in `kos/src/include/kostypes.h`.
- KOS assumes 4 cores only
- Only need the least significant 4 bits
- Examples:
 - Mask $0^{60}0001 \Rightarrow$ process can be only assigned to core 1
 - $0^{60}0101 \Rightarrow$ process can be only assigned to cores 1 & 3
 - $0^{60}1111 \Rightarrow$ process can be assigned to any of the 4 cores

Scheduler::preempt() in src/runtime/Scheduler.cc

```
void Scheduler::preempt() {           // IRQs disabled, lock count inflated
    #if TESTING_NEVER_MIGRATE
        switchThread(this);
    #else /* migration enabled */
        Scheduler* target = Runtime::getCurrThread()->getAffinity();    Replace for A1
    #if TESTING_ALWAYS_MIGRATE
        if (!target) target = partner;
    #else /* simple load balancing */
        if (!target) target = (partner->readyCount + 2 < readyCount) ? partner : this;
    #endif
        switchThread(target);
    #endif
}
```

Assignment 1

- Get the affinity mask
 - `mword affinityMask = Runtime::getCurrThread()->getAffinityMask()`
- If the affinity mask = 0, there should be no change in policy
 - `target = Runtime::getCurrThread()->getAffinity()`
 - Same as before; returns a scheduler object
- Else check every bit in affinity mask
 - If bit `i` is set, get the scheduler for that core
 - `Scheduler *sched = Machine::getScheduler(i)`
 - For all set bits, assign the process to the core with the smallest ready queue
 - Queue size: `sched->readyCount`
 - Assignment: `target = sched`